

Lecture 21 :

Dynamic Programming $\dim < 4 \text{ or } 5$, $c(), f(), x, u$ known

$$\hat{J}^*(x_i) \leftarrow \min_{u \in U} [c(x_i, u) + \hat{J}^*(f(x_i, u))]$$

$$\pi^*(x_i) = \operatorname{argmin}_{u \in U} [c(x_i, u) + J^*(f(x_i, u))]$$

Stability analysis via SOS $f()$ vector-valued poly
 $f(), x, u$, known

find s.t. $V_\alpha(x)$ is SOS
 α

$$-\dot{V}_\alpha(x) = -\frac{\partial V_\alpha}{\partial x} f(x) \text{ is SOS}$$

Trajectory Optimization $c(), f(), x, u$ known, local solutions

$$\begin{aligned} &\text{minimize} && \sum_{n=0}^{N-1} c(x_n, u_n) dt \\ &x_1, \dots, x_N, u_0, \dots, u_{N-1} \end{aligned}$$

$$\text{s.t. } x_{n+1} = x_n + f(x_n, u_n) dt$$

Pendulum DP

Airplane verification SOS

Rocket landing, walking / backflip robot traj opt

Flying a quadrotor in unknown env traj opt, LQR

A clarification

$$x = [x_{\text{robot}}, x_{\text{world}}]$$


Motion planning for cars inference other drivers motion

Sushi making ??? f, x, u, c

Markov decision process $M = \{X, U, f(\cdot), c(\cdot)\}$



Partially observed MDP $M = \{X, U, \gamma, f(\cdot), g(\cdot), c(\cdot)\}$



Estimate $f(\cdot)$

"system identification"

"model learning"

Estimate $g(\cdot)$, estimate x

"state estimation"

Estimate what X (the space of x) should even be
latent space learning?

Estimate $c(\cdot)$ from demonstrations of an expert

"Imitation Learning"

$$\begin{aligned} \min_{x_0, \dots, x_N, u_0, \dots, u_N} \quad & \sum_{i=0}^N c(x_i, u_i) \quad \text{non-linear programming} \\ \text{s.t.} \quad & x_{t+1} = f(x_t, u_t) \end{aligned}$$

Given not much known in manipulation

options:

1. find prob domains where most things are known
2. assume some subset is known
3. tackle hardest, general problem

known dynamics unknown dynamics unknown dynamics
known state known state partially observed state

→ →

"Copy-demo" MIT 1970

Ishikawa Komuro lab, 2009

Dense Object Nets visual object descriptors CORL 2018

one image → class of images description

key point affordances for category level manipulation

Imitation Learning

π_{expert}

↓

$D = \{T_0, T_1, \dots, T_m\}$

$T_i = ((y_0, u_0), (y_1, u_1), \dots, (y_{N_i}, u_{N_i}))$

π_{learner}

1. Behavior Cloning

Fit from data

$u = \pi(y_{-T}, y_{-T+1}, \dots, y)$

2. Inverse Optimal Control / Inverse Reinforcement Learning

a. infer $C()$

b. optimize π

3. Generative Adversarial Imitation Learning

Make π_{Learner}

indistinguishable

from π_{Expert}

Contact Rich Manipulation Tasks

Solution Approaches

Model Based:

controllers with good generalization
provides insides


- hard to hybrid, unknown dynamics

Learning based:

model-based RL

model-free RL

Model-based RL



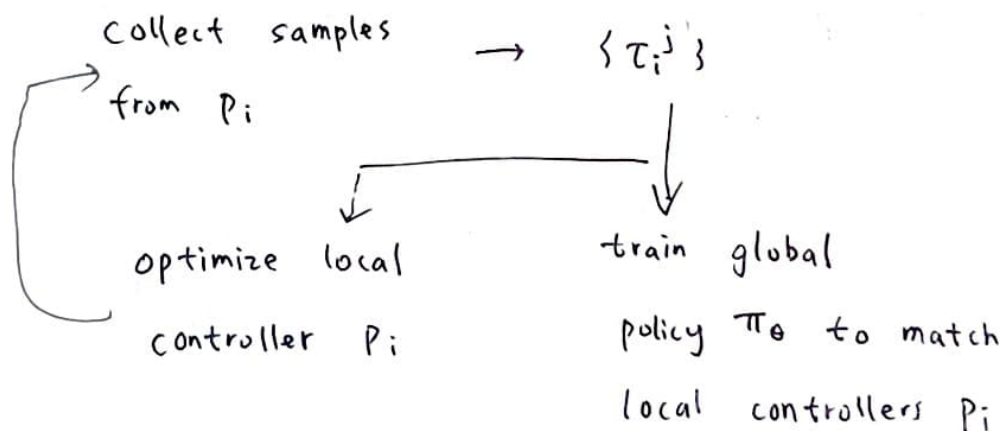
Motivation

Goal: controller for contact rich manipulation

Challenges: unknown dynamics
partially observable states

Overview

1. collect samples from robot $\{\tau_i^j\}$
2. fit time-varying linear-Gaussian dynamics
3. optimize local controllers using iLQG
4. Repeat 1-3 until convergence
5. train global policy by imitating local policies



$$\tau_j = \{x_i, u_i\}_{i=1}^{\tau_j}$$

iLQG with known dynamics

locally-optimal feedback control of nonlinear stochastic systems


$$\min_{\pi} E \left[h(x(\tau)) + \int_t^{\tau} \ell(\tau, x(\tau), \pi(\tau, x(\tau))) d\tau \right]$$

$$dx = f(x, u) dt + F(x, u) d\underline{w} \quad \text{random variable?}$$

i LQG

convert to
approach: \rightarrow LQG, something we know how to solve

1. current trajectory $\bar{x}(t), \bar{u}(t)$
2. approximate costs and trajectories around current trajectory
quadratic approximation of costs
linear approximation of dynamics
3. Solve linear quadratic problem to get local feedback controller $\delta u = I_k + L_k \delta x$
4. use controller to update current trajectory
5. Repeat 1-4 until convergence



Linear - Quadratic Approximation

Start with nominal trajectory $\bar{x}(t), \bar{u}(t)$

Discrete time dynamics $\bar{x}_{k+1} = \bar{x}_k + \Delta t \cdot f(\bar{x}_k, \bar{u}_k)$

Linearized Dynamics $\delta x_k \triangleq x_k - \bar{x}_k, \delta u_k = u_k - \bar{u}_k$

$$\delta x_{k+1} = A_k \delta x_k + B_k \delta u_k + E_k (\delta u_k) \xi_k$$

Quadratic approximation to cost

$$\begin{aligned} \text{cost}_k = & q_k + \delta x_k^T q_k + \frac{1}{2} \delta x_k^T Q_k \delta x_k \\ & + \delta u_k^T r_k + \frac{1}{2} \delta u_k^T R_k \delta u_k + \delta u_k^T P_k \delta x_k \end{aligned}$$

Solving single LQG step

Very similar to LQR

Cost-to-go function is quadratic under any affine control law

$$\delta u = I_k + L_k \delta x \quad V_k(\delta x) = S_k + \delta x^T S_k + \frac{1}{2} \delta x^T S_k \delta x$$

Can show via induction that

$$S_k = Q_k + A_k^T S_{k+1} A_k - G^T H^{-1} G$$

$$s_k = q_k + A_k^T s_{k+1} - G^T H^{-1} g$$

$$S_k = q_k + S_{k+1} + \frac{1}{2} \sum_i C_{i,k}^T S_{k+1} C_{i,k} - \frac{1}{2} g^T H^{-1} g$$

Extract optimal control law $\delta u = I_k + L_k \delta x$ by recursively minimizing cost to go

Bellman Equation

$$\begin{aligned} V_k(\delta x) &= \text{immediate cost} + E[V_{k+1}(\text{next state})] \\ &= q_k + \delta x^T (q_k + \frac{1}{2} Q_k \delta x) + \pi^T (r_k + \frac{1}{2} R_k \pi) \\ &\quad + \pi^T P_k \delta x + E[V_{k+1}(A_k \delta x + B_k \pi + C_k \xi_k)] \end{aligned}$$

$$\begin{aligned} V_k(\delta x) &= q_k + S_{k+1} + \frac{1}{2} \sum_i C_i^T S_{k+1} C_i + \delta x^T (q_k + A_k^T S_{k+1}) \\ &\quad + \frac{1}{2} \delta x^T (Q_k + A_k^T S_{k+1} A_k) \delta x \\ &\quad + \pi^T (g + G \delta x) + \frac{1}{2} \pi^T H \pi \end{aligned}$$

g, G, H depend only on known running cost, dynamics and cost-to-go at $k+1$

$$g \triangleq r_k + B_k^T s_{k+1} + \bar{\Sigma}_i C_{i,k}^T s_{k+1} C_{i,k}$$

$$G \triangleq P_k + B_k^T s_{k+1} A_k$$

$$H \triangleq R_k + B_k^T s_{k+1} B_k + \bar{\Sigma}_i C_{i,k}^T s_{k+1} C_{i,k}$$

Minimized control law $s_u = -H^{-1} (g + G \delta x)$

note other terms not related to π

But $H \neq 0$ may not be true, no control constraints

note s_u may be large step, need to be in approximate dynamics region

$$\nabla s_u = g + G \delta x$$

$$s_u = -\epsilon \hat{H}^{-1} (g + G \delta x)$$

$$\hat{H} > 0$$

can restrict s_u to a region, distance small

Peg-in-hole

Cost function: distance of known points on rigidly grasped object to their desired location

$$r_e (\|p_t - p_t^*\|)$$

State vector

joint angles and velocities

cartesian velocity on manipulated object

$$P_t - P_{\text{target}}$$

previously commanded torque

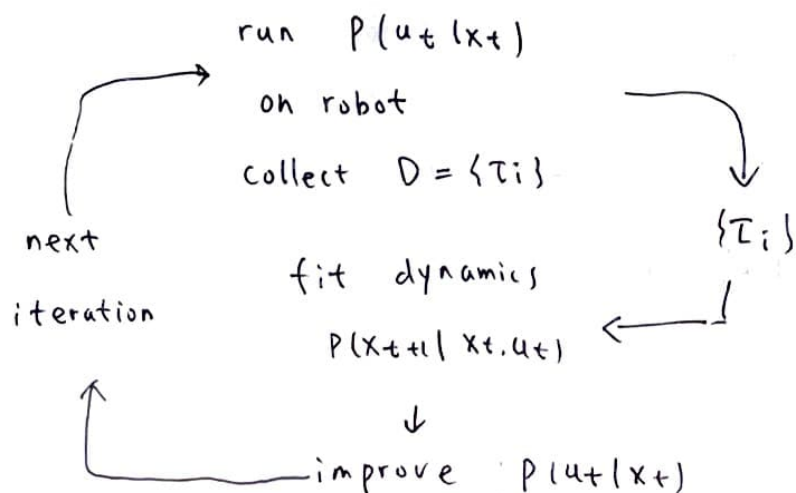
command = joint torque

iLQG with unknown dynamics

Problem statement: extend iLQG to settings where we don't know the dynamics

Approach:

1. Iteratively refit locally linear models to the dynamics
2. With locally linear models in hand run iLQG as before
3. Run step 1-2 until convergence



Fit locally - linear dynamics

$$p(x_{t+1} | x_t, u_t) = \mathcal{N}(A_t x_t + B_t u_t + c_t, \Sigma_t)$$
$$\{x_t, u_t, x_{t+1}\}$$

Need to have a good priori to make this step works

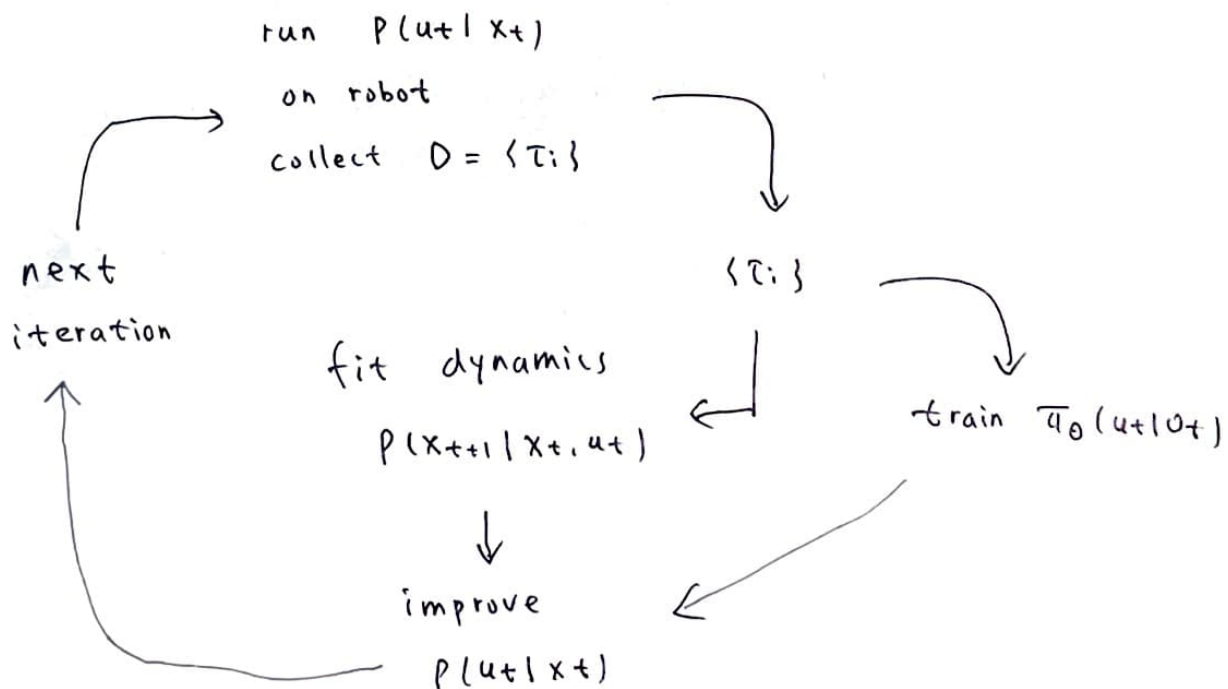
Time varying \rightarrow capture discontinuous dynamics?

average the ^{dynamics} ~~dynamics~~ from both sides of a discontinuity

problem: can only work in a region around trajectory



Guided Policy Search with Local Models





local policies

controller 1

controller 2

controller 3

train a global function approximator

$$\pi_{\theta}(x_t) = u_t$$

\uparrow can be image